

Long Latency Operations

Daniel Jiménez-González

March 6, 2018

Index

Index	i
1 Previous Work	1
1.1 Arithmetic Expression Optimizations	1
1.2 Memoization	1
2 Arithmetic Expression Optimizations	1
3 Memoization and Buffering	2
4 Routine Specialization	3

1

Previous Work

We provide you a *.zip* file with several input files.

- There are several Makefile targets to automatically generate the binary files. For instance, *make prog.pg* will generate an executable compiled with profiling flags. Using *CFLAGS* variable in the Makefile in a subdirectory (local Makefile) of a program you can define the specific compilation flags for a program. Using *PROGS* variable in the local Makefile you can do "make" of different program names. You can read README file in the root directory of the laboratory lab3.
- Most of the programs have a parameter in order to modify the amount of work to do. You can modify that parameter in order to obtain a reasonable execution time.
- Most of the programs use a random input data. However, as they don't modify the seed of the random function, they should obtain the same results for different executions.

Some advices for the lab activities:

- You must redirect the output of the program executions to a file since the output can be large and/or with not printable characters.
- To compare ASCII output files, you should use tools like *diff*, *tkdiff* or *cmp*.
- To compare binary output files, you should use *cmp* tool.
- In order to visualize a file with non printable characters you can use the *od* tool (octal dump). For instance, *od -b /bin/ls | more* shows the ASCII code of the characters of the file */bin/ls*. That tool may be parametrized with the starting point to do the **dump**.

In order to take profit of the activities at the laboratory you should do the following activities before going to the lab:

1.1 Arithmetic Expression Optimizations

1. We give you the **primers.c** program that uses the Eratostenes method in order to create a prime list (from 1 upto the number you indicate as parameter).
 - Look at <http://www.wikipedia.org> and the GCC low-level runtime library for information.
 - Try to find something regarding the `__udivdi3` and `__umoddi3`. That will help you with laboratory exercises. Hint: those symbols do not appear depending on the architecture you compile to.

1.2 Memoization

2. Look for information regarding to the trigonometric routine implementations `sin()` and `cos()`. Is there any hardware instruction to do it in the Cortex A9 of our boards?

2

Arithmetic Expression Optimizations

1. `primers.c` program uses the Eratostenes method in order to create a prime list (from 1 upto the number you indicate as parameter).
 - (a) Compile `primers` with `"-O0"` (`make primers.0` and do timing, and keep the original output result.
 - (b) Use profiling (with `gprof`, `valgrind` and/or `operf`, `opannotate`) to analyze its behavior when compiling with `-O0`. Note: you may want to use the `primers` argument to reduce the computation time with `valgrind`.
 - (c) Which are the most time consuming functions? Look at their most time consuming lines to figure out which are the expensive operations. Is there any way to reduce/avoid those costly operations?
 - (d) Copy the original code to `primers_opt.c`. Optimize this copy of `primers` based on the analysis done in the previous question.
 - (e) Compile it using `make primers_opt.g` and check that the results of the optimized version you have done are the same than the original one.
 - (f) Which speedup have you obtained with your optimized code compared to the original `primers` when compiling both of them with `"-O0"`?
 - (g) Now, compile the original and the optimized versions of `primers` with `"-O3"` and do timing of both of them.
 - (h) Which speedup have you obtained compared to the original `primers` when compiling with `"-O3"` both programs, original and optimized? Do profiling again and look at the output of `objdump -d` in order to explain the speedup obtained.

3

Memoization and Buffering

2. `trigon.c` program performs several calls to write and the trigonometric routines `sin()` and `cos()`.
 - (a) Study the source code of the program.
 - (b) Compile the program with `make trigon.pg3`. Profile and analyze the program behavior. Explain the results.
 - (c) Re-compile the program with `make trigon.pg3s`. Profile and analyze the program behavior. Explain the results.
 - (d) Figure out how many system calls are done in the program, and which is the elapsed time of the program. Note that `strace` may take a while to find out this information.
 - (e) Based on the previous analysis, modify the program in order to avoid so many calls and reduce the overall elapsed time.
 - (f) Compute the speed-up of your optimized version compared to the original version.
 - (g) Profile the new code. How much CPU time is devoted to trigonometric computations? Which is the maximum speedup that we can achieve if we improve the corresponding execution time?
 - (h) Modify the program in order to avoid the computation repetitions and check that the output of the new version is the same as the original code.
 - (i) Profile the new version of the program and compare it with the profile of the original code.
 - (j) Which overall speed-up have you obtained?

4

Routine Specialization

3. `pi.c` program compute the first 10000 decimals of the π number.

- (a) What is the speedup of the `pi` program compiled with `O3` compared to the `pi` compiled with `O0`?
- (b) Profile both `pi` binaries (compiled with `O3` and with `O0`) using `valgrind` and `callgrind.annotate`. Open the assembler code to understand the differences (`objdump -d`) and/or look for the places where "divide instructions" are executed using `callgrind.annotate` (Note: use 1000 as argument of `pi.c` when using `valgrind`). Could you justify these differences?
- (c) It seems that the compiler has done a good work but there is still work to do. Based on the previous profiling, look at the source code of the program, and in particular, the source code of routine `DIVIDE`. Then, propose an optimization that can help to reduce the cost of each `specific` call to `DIVIDE` in order to reduce/avoid long latency operations. Specilize the code of `pi` using memoization. Some hints:
 - For each `DIVIDE` by a particular value, `x[k]`, `r`, and `u` in the code will always have the same range of values (maybe different range among them)
 - Those values in the range are for a limited number of input values
- (d) Profile the new version of the program and compare it with the profile of the original code.
- (e) Which speed-up have you obtained?